


The N-Body Problem Optimized with MPI



John Sabo


What is the N-Body problem?

- “In **physics**, the **n -body problem** is the problem of predicting the individual motions of a group of **celestial objects** interacting with each other **gravitationally**.” [1]



Inspiration for Project

Solar System Simulation



Objective

- Implement main calculation in C
- Optimize program with OpenMPI
- Verify identical output
- Test performance



Serial Algorithm (Javascript)

- Runtime : $O(n^2)$
- Each body takes into account every other body's position and velocity in order to determine its next position and velocity.

```
for(var planetA of planets)
  for(var planetB of planets)
    if(planetA != planetB)
      var denominator = Math.pow(
        Math.pow(planetA.x - planetB.x, 2)
        + Math.pow(planetA.y - planetB.y, 2)
        + Math.pow(planetA.z - planetB.z, 2), 3/2.0)

      nextVX += (G * planetB.mass * (planetA.x - planetB.x)
        * parameters.timeStep) / denominator
      nextVY += (G * planetB.mass * (planetA.y - planetB.y)
        * parameters.timeStep) / denominator
      nextVZ += (G * planetB.mass * (planetA.z - planetB.z)
        * parameters.timeStep) / denominator
```

$$x'_j = vx_j$$

$$y'_j = vy_j$$

$$z'_j = vz_j$$

$$vx'_j = \frac{G \sum_{j \neq k} m_k (x_k - x_j)}{[(x_j - x_k)^2 + (y_j - y_k)^2 + (z_j - z_k)^2]^{\frac{3}{2}}}$$

$$vy'_j = \frac{G \sum_{j \neq k} m_k (y_k - y_j)}{[(x_j - x_k)^2 + (y_j - y_k)^2 + (z_j - z_k)^2]^{\frac{3}{2}}}$$

$$vz'_j = \frac{G \sum_{j \neq k} m_k (z_k - z_j)}{[(x_j - x_k)^2 + (y_j - y_k)^2 + (z_j - z_k)^2]^{\frac{3}{2}}}$$

MPI vs Sequential Data Handling

Sequential:

```
all_Bodies = [ Body1: { x, y, z, vx, vy, vz, m },  
               Body2: { x, y, z, vx, vy, vz, m },  
               Body3: { x, y, z, vx, vy, vz, m },  
               Body4: { x, y, z, vx, vy, vz, m },  
               Body5: { x, y, z, vx, vy, vz, m },  
               Body6: { x, y, z, vx, vy, vz, m },  
               Body7: { x, y, z, vx, vy, vz, m },  
               ...  
               ...  
               ...  
               BodyN: { x, y, z, vx, vy, vz, m } ]
```

MPI:

0
all_Bodies, local_B = all_Bodies[0:A]

1
all_Bodies, local_B = all_Bodies[A+1:C]

...

N
all_Bodies, local_B = all_Bodies[C+1:N]

Serial int main()

1. Allocate Memory for **bodies**
2. Initialize each **body** with random values
3. Run the calculation for **K** steps
4. Print time taken
5. Free memory

MPI int main()

1. Allocate memory for **all_bodies** and **local_bodies**
2. Based on **N** establish **local_body** count and **displacement** in **all_bodies** for each process. (This allows data from each process to be sent to the correct index in **all_bodies** in every other process)
3. Initialize **all_bodies** within process of **rank zero**.
4. Broadcast **all_bodies** from **rank zero** to **all** other processes.
5. ScatterV **all_bodies** to every process evenly into **local_bodies**.
6. Run the calculation for **K** steps on **local_bodies** in each process.
7. After each calculation, use All-GatherV to update **all_bodies** with each processes' updated **local_bodies**. Utilize **displacement** to ensure **local_bodies** are entered into correct index within **all_bodies**.
8. Print time taken from **rank zero** and free memory.

1) Allocate memory for *all_bodies* and *local_bodies*

```
Body *localBodies = malloc(local_n * sizeof(Body));  
Body *allBodies = malloc(numBodies * sizeof(Body));
```

2) Establish displacement indexes for each process

```
int *recvcounts = malloc(world_size * sizeof(int));
int *displs = malloc(world_size * sizeof(int));
int sum = 0;
for (int i = 0; i < world_size; i++) {
    // data amount for process i
    recvcounts[i] = (numBodies / world_size) * sizeof(Body);
    // starting index in allBodies for process i
    displs[i] = sum;
    sum += recvcounts[i];
}
```

3) Initialize All_Bodies

```
if (world_rank == 0) {  
    initializeBodies(allBodies, numBodies);  
}
```

4) Broadcast All_Bodies to all processes

```
MPI_Bcast(allBodies, numBodies * sizeof(Body), MPI_BYTE, 0, MPI_COMM_WORLD);
```

`allBodies` -> receiving/sending buffer

`numBodies * sizeof(Body)` -> data size

`MPI_BYTE` -> data type

`0` -> root process

`MPI_COMM_WORLD` -> communicator

5) ScatterV All_Bodies to each processes' Local_Bodies

```
MPI_Scatterv(allBodies, recvcounts, displs, MPI_BYTE, localBodies, local_n *  
sizeof(Body), MPI_BYTE, 0, MPI_COMM_WORLD);
```

allBodies	-> sending buffer
recvcounts	-> sent data size
displs	-> send from index
MPI_BYTE	-> sent data type
localBodies	-> receiving buffer
local_n * sizeof(Body)	-> received data size
MPI_BYTE	-> received data type
0	-> root process
MPI_COMM_WORLD	-> communicator

6/7) Run the simulation and All-GatherV

```
for (int step = 0; step < numSteps; step++) {  
    // calculation same as sequential except outer loop uses localBodies  
    computeForcesAndMove (localBodies, allBodies, local_n, numBodies);  
    MPI_Allgatherv (localBodies, local_n * sizeof(Body), MPI_BYTE, allBodies,  
recvcounts, displs, MPI_BYTE, MPI_COMM_WORLD);  
}
```

localBodies	-> sending buffer
local_n * sizeof(Body)	-> sent data size
MPI_BYTE	-> sent data type
allBodies	-> receiving buffer
recvcounts	-> received data size
displs	-> receive into index
MPI_BYTE	-> received data type
MPI_COMM_WORLD	-> communicator

8) Print time taken and free memory

```
if(world_rank == 0)
{
    printf("%lf\n", endtime - starttime);
}
```

```
free(localBodies);
free(allBodies);
free(recvcounts);
free(displs);
MPI_Finalize();
```

Verification

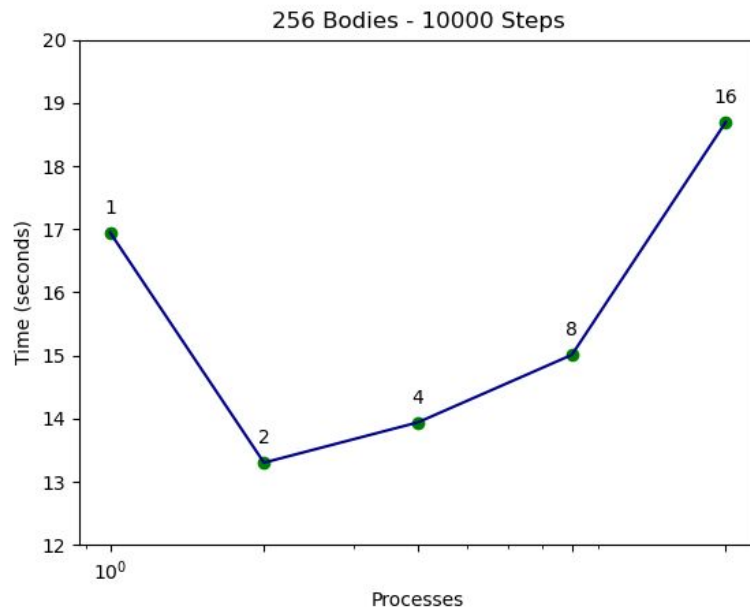

- Run sequential and MPI scripts with set random seeds and equal body count / step count.
- Print each body's final position and velocity into a file.
- Utilize diff command to check output files are identical.

Outputs are identical for 10 bodies and 10 generations.

Outputs are identical for 100 bodies and 100 generations.

Outputs are identical for 1000 bodies and 1000 generations.

Performance

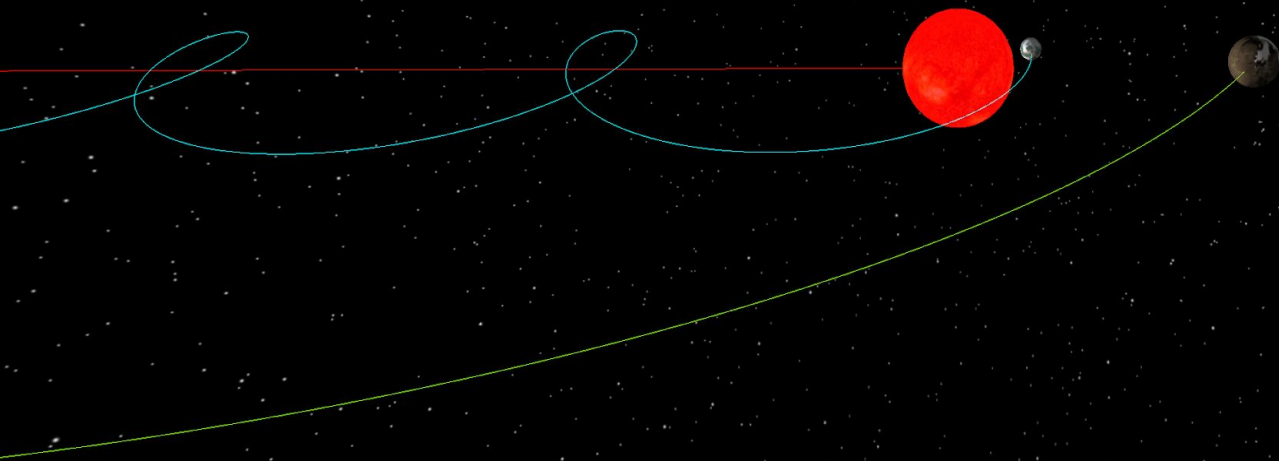


Analysis and Conclusion

- **Speed up decreases** with process count **greater than two** due to **overhead** associated with communicating with every process each step in the simulation.
- Testing on a machine with **more than two cores** may demonstrate **greater speed up**.
- Successful implementation of **OpenMPI!**

What's next?

- Implement visualization of the simulation.
- Optimize with CUDA?
- How can we avoid communicating excess information between processes?



:(**Thanks!**

Your PC ran into a problem and needs to restart. We're just collecting some error info, and then we'll restart for you.


20% complete



For more information about this issue and possible fixes, visit <https://www.windows.com/stopcode>

If you call a support person, give them this info:

Stop code: CRITICAL_PROCESS_DIED



Sources

- [1] https://en.wikipedia.org/wiki/N-body_problem
- [2] <https://www.open-mpi.org/doc/>
- [3] <https://www.mpich.org/static/docs/latest/www3/>